


# Hypothesis Testing with Randomization



# Housekeeping

- Office hours change this week

# Where we're going today

- We will see another kinds of hypotheses for different types of research questions
- Hypothesis testing framework is the same, but will change how we obtain null distribution
- Try to see the big picture

# Test of independence

# Running example: sex discrimination study

- Note: this study considered sex as binary “male” or “female”, and did not take into consideration gender identities
- Participants in the study were 48 bank supervisors who identified as male and were attending a management institute at UNC in 1972
  - Each supervisor was asked to assume the role of personnel director of a bank
  - They were each given a file to judge whether the person in the file should be promoted
  - The files were identical, except half of them indicated that the candidate was male, and the other half were indicated as female
  - Files were randomly assigned to bank managers
  - Experiment or observational study?
- Research question: Are individuals who identify their sex as female discriminated against in promotion decisions made by their managers who identify as male?

# Defining hypotheses

Research question: Are individuals who identify their sex as female discriminated against in promotion decisions made by their managers who identify as male?

- What is/are the variables(s) here? What types of variables are they?
- We need to construct hypotheses where  $H_0$  is “status quo” and  $H_A$  is the claim researchers have
- $H_0$ : the variables **sex** and **decision** are independent.
  - i.e. any observed difference in promotion rates is due to variability
- $H_A$ : the variables **sex** and **decision** are *not* independent, and equally-qualified female personnel are less likely to be promoted than male personnel

# Data

For each of the 48 supervisors, the following were recorded:

- The sex of the candidate in the file (male/female)
- The decision (promote/not promote)

sex	not promote	promote	total
female	10	14	24
male	3	21	24
total	13	35	48

- Are we prepared to answer our research question: Are individuals who identify their sex as female discriminated against in promotion decisions made by their managers who identify as male?
  - What evidence do we have?

# Data (cont.)

Conditional probability of getting promoted by sex:

```
1 discrimination |>
2   count(sex, decision) |>
3   group_by(sex) |>
4   mutate(cond_prob = n/sum(n)) |>
5   filter(decision == "promote") |>
6   select(-n)
```

```
# A tibble: 2 × 3
# Groups:   sex [2]
  sex    decision cond_prob
<chr>  <chr>      <dbl>
1 female promote    0.583
2 male   promote    0.875
```

- Is the observed difference -0.2916667 convincing evidence? We need to examine variability in the data, assuming  $H_0$  true.
- Let's set  $\alpha = 0.05$



# Simulate under null

- Simulating under  $H_0$  means operating in a hypothetical world where **sex** and **decision** are independent.
  - This means that knowing the **sex** of the candidate should have no bearing on the **decision** to promote or not
- We will perform a simulation called a **randomization test** where we randomly re-assign **decision** and **sex** outcome pairs to see what would have happened if the bankers' **decision** had been independent of candidate's **sex** (i.e. if  $H_0$  true)

# Randomization test

sex	not promote	promote	total
female	10	14	24
male	3	21	24
total	13	35	48

- Write down “promote” on 35 cards and “not promote” on 13 cards. Repeat the following:
  - Thoroughly shuffle these 48 cards.
  - Deal out a stack of 24 cards to represent males, and the remaining 24 cards to represent females
    - This is how we simulate independence under  $H_0$
  - Calculate the proportion of “promote” cards in each stack,  $\hat{p}_{m,sim}$  and  $\hat{p}_{f,sim}$
  - Calculate and record the difference  $\hat{p}_{f,sim} - \hat{p}_{m,sim}$  (order of difference doesn't matter so long as you are consistent)

# Randomization test (activity)

Try it!

# Randomization test (code)

Let's perform one iteration of the simulation.

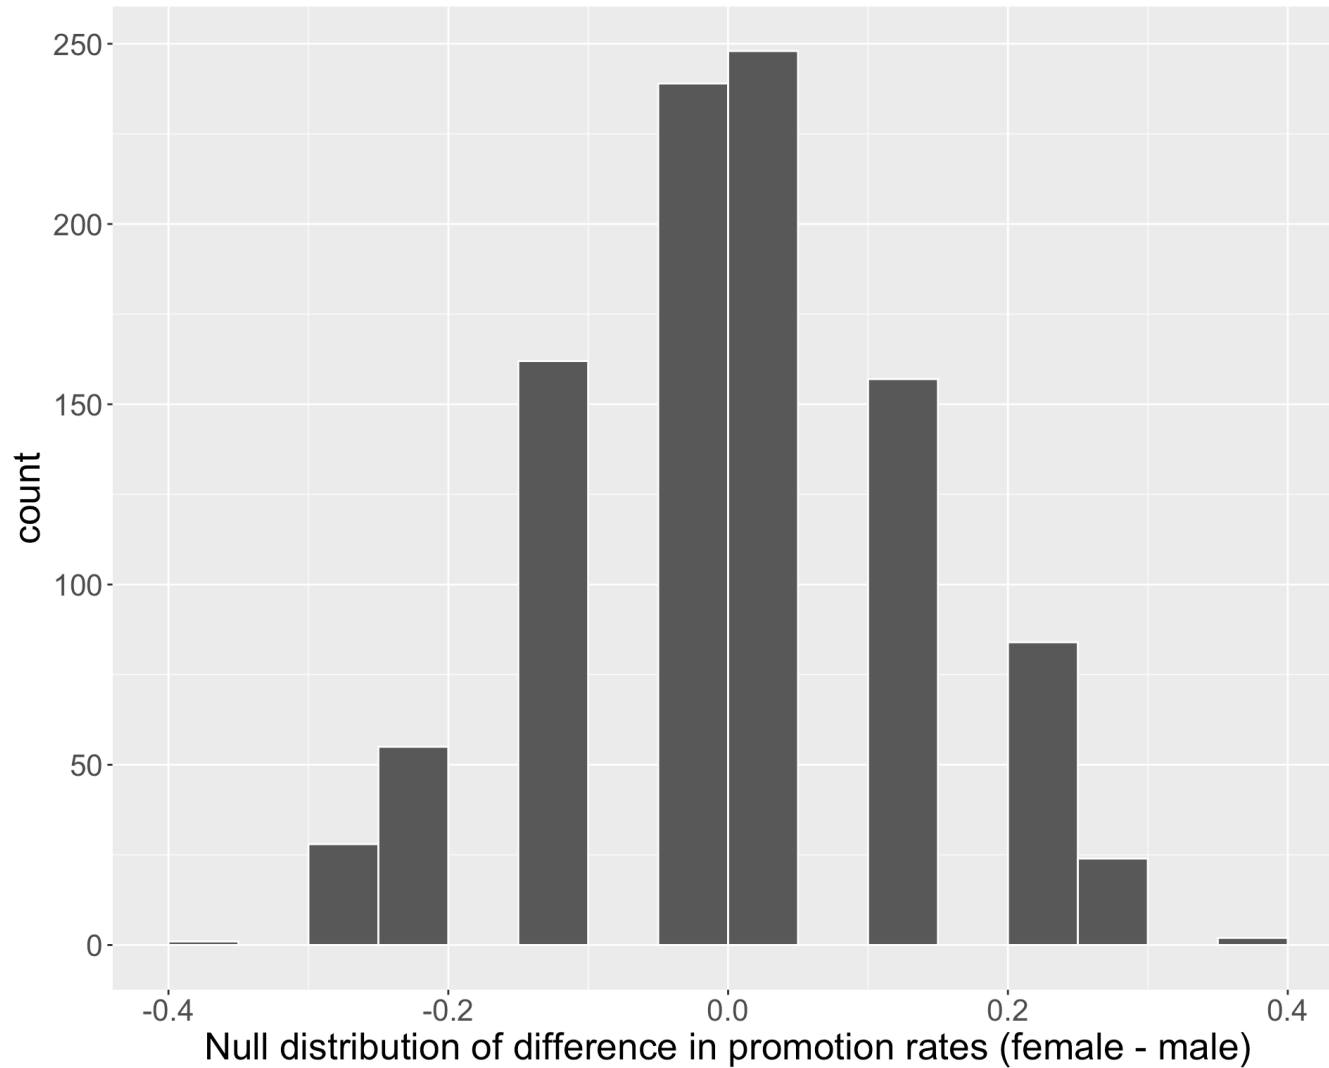
```
1 # reproducibility
2 set.seed(1)
3 n_female <- sum(discrimination$sex == "female")
4 n_male <- sum(discrimination$sex == "male")
5
6 # create cards
7 cards <- discrimination$decision
8
9 # shuffle cards
10 shuffled <- sample(cards)
```

sex	not promote	promote	total
female	8	16	24
male	5	19	24
total	13	35	48

- Under this simulation, 0.6666667 of females were promoted, and 0.7916667 of males were promoted. Simulated difference: -0.125

# Null distribution

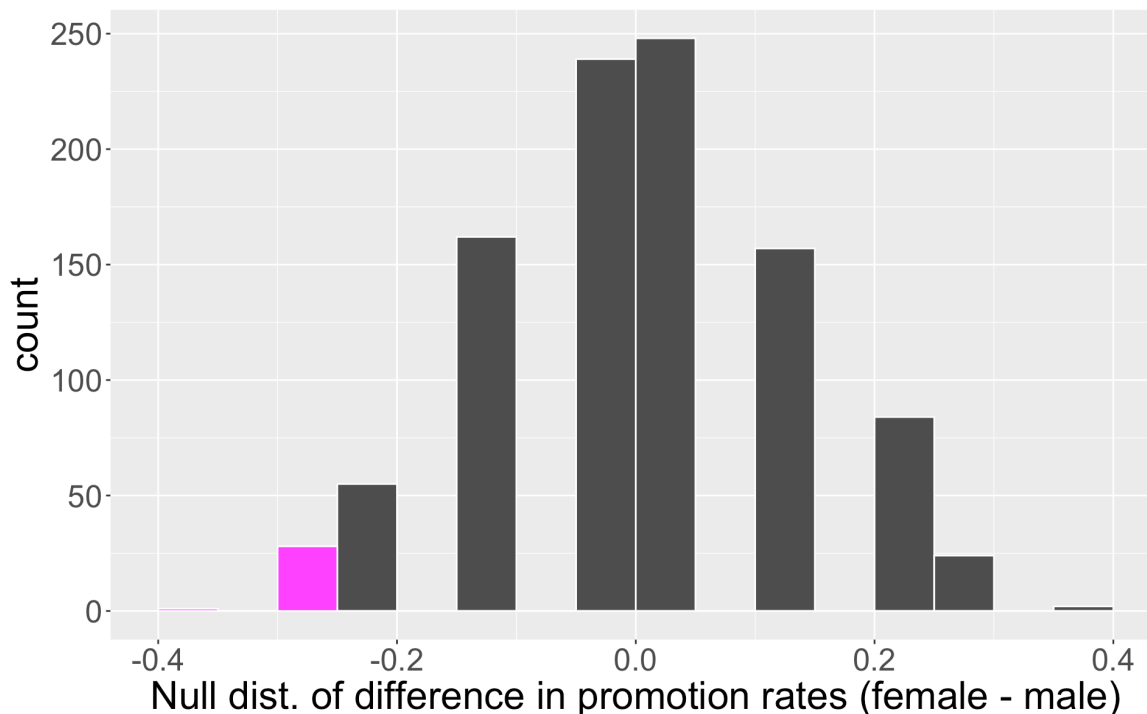
Repeat the previous step 1000 times:



# Obtain p-value

Recall, the observed difference in our data was  $\hat{p}_f - \hat{p}_m = -0.2916667$ .

- p-value is probability of observing data as or more extreme than our original data, given  $H_0$  true.
- Where does “as or more extreme” correspond to on our plot?



- Out of 1000 simulations under  $H_0$ , 29 resulted in a difference in promotion rates as or more extreme than our observed
- So the p-value is 0.029



# Making decision and conclusion

Our research question: Are individuals who identify their sex as female discriminated against in promotion decisions made by their managers who identify as male?

- $H_0$ : sex and decision are independent
  - $H_A$ : sex and decision are not independent and equally-qualified female personnel are less likely to get promoted than male personnel by male supervisors
  - $\alpha = 0.05$
- 
- Interpret our p-value in context.
  - Make a decision and conclusion in response to the research question.



# Making decision and conclusion (answer)

- Assuming that **sex** and **decision** are independent, the probability of observing a difference in promotion rates as or more extreme as we did is 0.029.
- Because the observed p-value of 0.029 is less than our significant level 0.05, we reject  $H_0$ . The data provide strong evidence of sex discrimination against female candidates by the male supervisors.
- What kind of error could we have made?

# Comparing two proportions

# Running example: CPR

An experiment was conducted, consisting of two treatments on 90 patients who underwent CPR for a heart attack and subsequently went to the hospital. Each patient was randomly assigned to either:

- treatment group: received a blood thinner
- control group: did not receive a blood thinner

For each patient, the outcome recorded was whether they survived for at least 24 hours.

group	died	survived	total
control	39	11	50
treatment	26	14	40
total	65	25	90

- What is/are the variables(s) here? What types of variables are they?

# Defining hypotheses

The researchers are interested in learning if the blood thinner treatment was effective.

In words, try to determine  $H_0$  and  $H_A$ .

- Let  $p_T$  and  $p_C$  denote the proportion of patients who survive when receiving the thinner (Treatment) and when not receiving the treatment (Control), respectively

Option 1

- $H_0: p_T \leq p_C$
- $H_A: p_T > p_C$

Option 2 (preferred)

- $H_0: p_T - p_C \leq 0$
- $H_A: p_T - p_C > 0$

# Collect data

Using the data, obtain the observed difference in sample proportions.

group	died	survived	total
control	39	11	50
treatment	26	14	40
total	65	25	90

```
1 p_hat_c <- cpr |>
2   filter(group == "control") |>
3   summarise(p = mean(outcome == "survived")) |>
4   pull()
5 p_hat_t <- cpr |>
6   filter(group == "treatment") |>
7   summarise(p = mean(outcome == "survived")) |>
8   pull()
9 obs_diff <- p_hat_t - p_hat_c
```

- $\hat{p}_C = \frac{11}{50} = 0.22$
- $\hat{p}_T = \frac{14}{40} = 0.35$
- Observed difference:  
 $\hat{p}_T - \hat{p}_C = 0.13$

- Is this “convincing evidence” that blood thinner usage after CPR is effective?
- Set  $\alpha = 0.05$

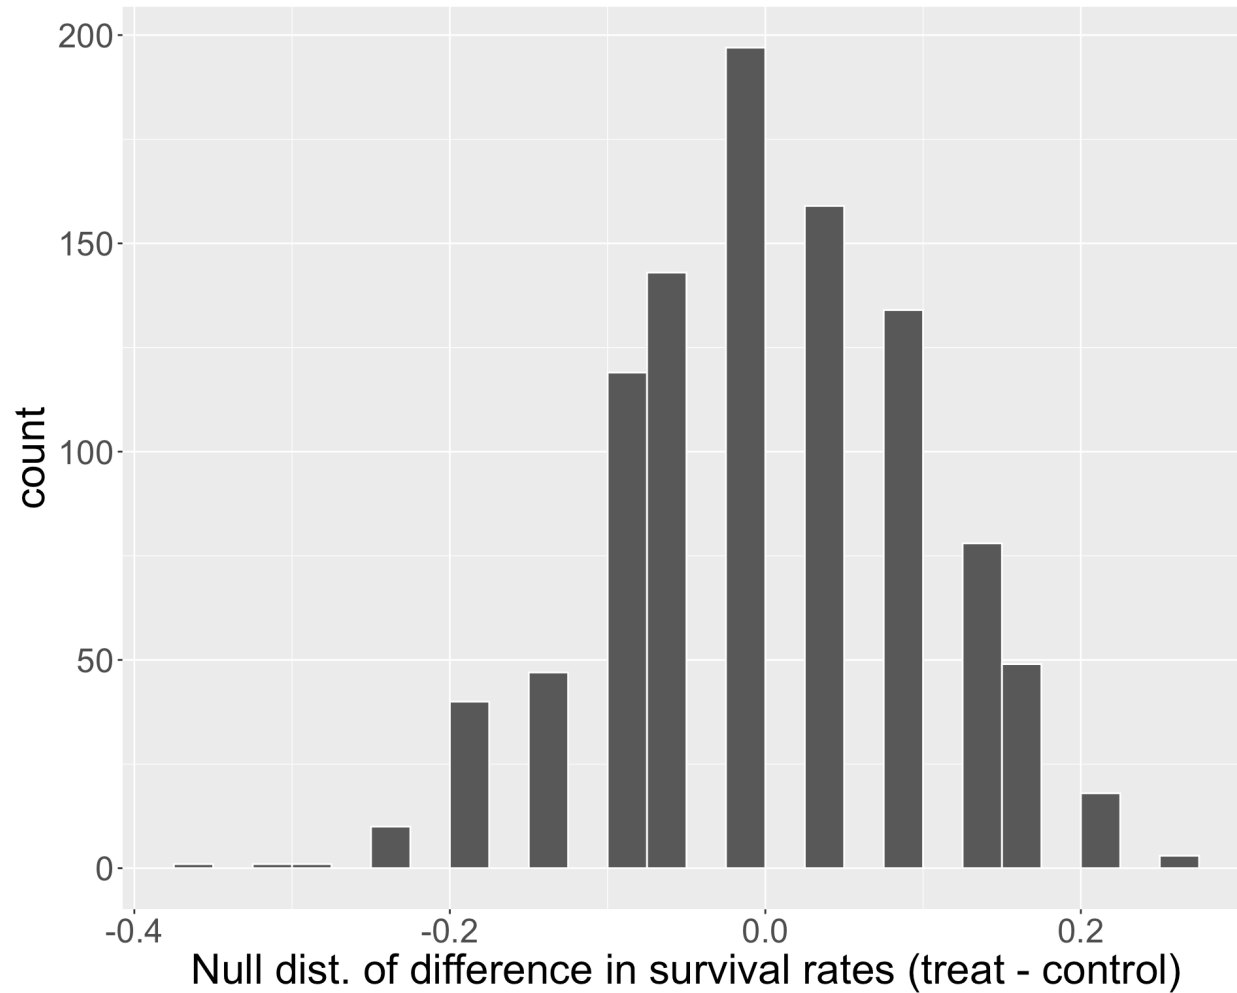
# Simulate under null

- We will once again perform a *randomization* test to try and simulate the difference in proportions under  $H_0$ 
  - Under  $H_0$ , treatment group is no better than control group, so let's simulate assuming that outcome and treatment are independent
- Write down **died** on 65 cards, and **survived** on 25 cards. Then repeat several times:
  - Shuffle cards well
  - Deal out 50 to be Control group, and remaining 40 to be Treatment group
  - Calculate proportions of survival  $\hat{p}_{C,sim}$  and  $\hat{p}_{T,sim}$
  - Obtain and record the simulated difference  $\hat{p}_{T,sim} - \hat{p}_{C,sim}$

# Simulate under null (code)

Live code


# Visualizing null distribution



How would we obtain the p-value in this problem?



# Calculate p-value



- We observed 148 out of 1000 simulations where the difference in proportions under  $H_0$  was as or more extreme than our observed difference of 0.13
- So p-value is 0.148

# Interpret and make conclusion

The researchers are interested in learning if the blood thinner treatment was effective.

Our p-value is 0.148 and our selected significance level was  $\alpha = 0.05$ .

- Make a decision and conclusion about the research question in context.

# Comprehension questions

- What were the similarities and differences between:
  - hypothesis test for independence
  - hypothesis test for two proportions
- How do the randomization tests today differ from the test for one proportion that we learned last class?

